

The apply() family

Daniela Castro-Camilo

General use

- **apply(x, index, function):** Apply a function to the rows (index=1) or columns (index=2) of a matrix.

```
set.seed(4032021)
A = matrix(sample(1:20, 15), 5, 3); A
```

```
##      [,1] [,2] [,3]
## [1,]  12    6   20
## [2,]  11    4   18
## [3,]   8   13    7
## [4,]   9    3   10
## [5,]   1    5   19
```

```
f = function(x) (max(x) + min(x))/2
apply(A, 2, f)
```

```
## [1]  6.5  8.0 13.5
```

```
apply(A, 1, f)
```

```
## [1] 13.0 11.0 10.0  6.5 10.0
```

- **lapply(x, function):** Apply a function to each element of the list x.

```
x = list(c(1,2,3), c(4,5,6)); x
```

```
## [[1]]
## [1] 1 2 3
##
## [[2]]
## [1] 4 5 6
```

```
lapply(x, f)
```

```
## [[1]]
## [1] 2
##
## [[2]]
## [1] 5
```

```
x = list("x1" = A, "x2" = A+2); x
```

```
## $x1
##      [,1] [,2] [,3]
## [1,]  12    6   20
## [2,]  11    4   18
## [3,]   8   13    7
```

```
## [4,]    9    3   10
## [5,]    1    5   19
##
## $x2
##      [,1] [,2] [,3]
## [1,]   14    8   22
## [2,]   13    6   20
## [3,]   10   15    9
## [4,]   11    5   12
## [5,]    3    7   21
```

```
g = function(y) apply(y, 2, f)
lapply(x, g)
```

```
## $x1
## [1]  6.5  8.0 13.5
##
## $x2
## [1]  8.5 10.0 15.5
```

```
lapply(x, sum) # output is a list
```

```
## $x1
## [1] 146
##
## $x2
## [1] 176
```

- **sapply(x, function):** Same as lapply. The difference is that sapply will try to simplify as much as possible the output of lapply.

```
sapply(x, sum) # output is a vector (simplified list)
```

```
## x1 x2
## 146 176
```

```
set.seed(4032021)
temp = round(runif(10, 0, 20),1)
h = function(x,y) {mean(x) > y}
sapply(temp, h, y = 15) # output is logical
```

```
## [1] FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE
```

- **vapply(x, function):** Same as sapply. The difference is that vapply allows you to specify the type of output.

```
vapply(temp, h, numeric(1), y = 15) # output is numeric (1 = TRUE, 0 = FALSE)
```

```
## [1] 0 0 1 0 1 0 0 0 0 0
```

- **tapply(x, y, function):** Apply a function to subsets of a vector x. The subsets are defined by the vector y.

```
x = 1:10
y = rep(c(T,F), 5)
tapply(x, y, sum)
```

```
## FALSE TRUE
##    30   25
```

- **mapply(function, x, y, ...)**: Apply a function on multiple objects by elements (the **m** stands for multivariate).

```
Q1 <- matrix(c(rep(1, 4), rep(2, 4), rep(3, 4), rep(4, 4)),4,4); Q1
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
## [2,]    1    2    3    4
## [3,]    1    2    3    4
## [4,]    1    2    3    4
```

```
# Faster using mapply()
Q2 <- mapply(rep,1:4,4); Q2
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
## [2,]    1    2    3    4
## [3,]    1    2    3    4
## [4,]    1    2    3    4
```

- Parallelised versions of **lapply** are available in, e.g., the package **parallel**. These are useful when parallelising a code that needs to run on several cores. If interested, look for **mclapply**, **mcmapply**.

Data example

I illustrate the use of some of the functions in the **apply()** family using the dataset **airquality** available in R.

```
data(airquality)
head(airquality)
```

```
##   Ozone Solar.R Wind Temp Month Day
## 1    41     190   7.4   67     5   1
## 2    36     118   8.0   72     5   2
## 3    12     149  12.6   74     5   3
## 4    18     313  11.5   62     5   4
## 5    NA      NA  14.3   56     5   5
## 6    28      NA  14.9   66     5   6
```

```
?airquality
```

```
# Data contains information on mean ozone, solar radiation, average wind speed
# and maximum daily temperature
```

1. **[Apply]** Find the days with the highest maximum temperature and mean ozone.

```
# `id` contains the row indices of `irquality`
# with the highest ozone and temperate
id = apply(airquality[, c("Ozone","Temp")], 2, which.max)
airquality$Day[id[1]] # day with highest ozone
```

```
## [1] 25
```

```
airquality$Day[id[2]] # day with highest temperate
```

```
## [1] 28
```

2. Find the number of days where solar radiation exceeded 180

```
which(airquality$Solar.R > 180)
```

```
## [1] 1 4 7 10 12 13 14 16 17 19 22 26 29 30 31 32 33
## [18] 34 35 36 37 39 40 41 42 43 45 46 47 48 55 62 63 64
## [35] 67 68 69 70 73 75 77 78 79 80 81 83 84 85 86 89 90
## [52] 91 92 99 100 101 102 104 105 111 112 113 115 116 117 118 120 121
## [69] 122 123 125 126 127 130 131 132 133 134 135 136 139 140 142 143 144
## [86] 149 151 153
```

3. `[lapply/sapply]` Repeat 2 for the sequence of thresholds `thres = seq(180, 200)`.

```
thres = seq(180, 200)
```

```
# Using for()
out1 = list()
for(i in 1:length(thres)){
  out1[[i]] = which(airquality$Solar.R > thres[i])
}
```

```
# Using lapply
f = function(th,x) {which(x > th)}
out2 = lapply(thres, f, x = airquality$Solar.R)
```

`lapply` returns a list, which in the case is convenient since the length of the output changes with the threshold:

```
length(out2[[1]])
```

```
## [1] 88
```

```
length(out2[[20]])
```

```
## [1] 75
```

`sapply` is useful as well, but the output is exactly the same as with `lapply`. This is because `sapply` can't simplify a list whose elements have different lengths.

```
# Using sapply
out3 = sapply(thres, f, x = airquality$Solar.R)
```

3. `[lapply]` `lapply` can be very useful for data processing. Say we have 100 datasets measuring the same information contained in `airquality`. These 100 datasets are contained in a list called `my.airquality`. For the sake of this example, I will create `my.airquality` by randomly sampling from `airquality` 100 times:

```
# Creating `my.airquality`
my.airquality = list()
for(i in 1:100){
  id = sample(1:nrow(airquality), nrow(airquality), replace = T)
  my.airquality[[i]] = airquality[id, ]
  rownames(my.airquality[[i]]) = c() # set rownames to be 1,2,...,nrow(airquality)
}
# check the first entry
head(my.airquality[[1]])
```

```
## Ozone Solar.R Wind Temp Month Day
```

```
## 1    45    212  9.7   79    8   24
## 2    13     27 10.3   76    9   18
## 3    48    260  6.9   81    7   16
## 4    40    314 10.9   83    7    6
## 5    30    193  6.9   70    9   26
## 6     7     49 10.3   69    9   24
```

Assume now that you want to create a new data frame that contains the solar radiation of all the 100 datasets as columns. Your new data frame, which we will call `all.solar`, should have `nrow(airquality)` rows and 100 columns. The following line of code selects the column `Solar.R` from the first element in `my.airquality`.

```
tmp = my.airquality[[1]][ "Solar.R"]
```

If we want to select the column `Solar.R` from all the element in `my.airquality`, we can use `lapply`:

```
all.solar = lapply(my.airquality, function(x) x[, "Solar.R"])
class(all.solar)
```

```
## [1] "list"
```

Note that `all.solar` is a list, not a data frame. To convert `all.solar` to a data frame, we can do:

```
all.solar = data.frame(matrix(unlist(all.solar), ncol = length(all.solar), byrow = FALSE))
names(all.solar) = paste0('dat', 1:100)
```